Training Corpus Matters: Byte-Level BPE Compression on Cloud and Technical Documentation

A controlled, byte-matched case study measuring tokenizer compression, not model quality

⟨AUTHOR NAME⟩*
⟨AFFILIATION⟩

2026

Abstract

General-purpose tokenizers may allocate vocabulary inefficiently for specialized technical documentation, inflating sequence length. We present a controlled, CPU-only case study of how the *training corpus* changes byte-level BPE tokenizer *compression* on unseen documentation, holding the algorithm, pre-tokenizer, decoder, normalization, special tokens, vocabulary size, and training-byte budget fixed. Under a byte-matched 32 MiB budget, we train General, Azure-documentation, and 50/50 Mixed tokenizers at 8K/16K/32K vocabularies and evaluate on nine leakage-controlled held-out sets. On an unseen cross-service Azure holdout (five whole held-out services, 964 documents), the Azure tokenizer uses 36.68% fewer tokens per 1,000 UTF-8 bytes than the byte-matched General tokenizer at 16K (95% paired-bootstrap CI $[35.77\%, 37.62\%]$; document win rate 1.000). Full specialization carries a general-English regression of -34.75%; a 50/50 Mixed tokenizer retains $+34.21\%$ of the Azure gain with only -4.39% general-English regression. The advantage is stable across vocabulary sizes, persists on prose-dominant content, and transfers partially to other Microsoft and external (Kubernetes) technical documentation. We report *compression only*: no model-quality, inference-latency, GPU-throughput, or cost claim is made. Code, source manifests with pinned commits, per-document result rows, and a machine-verified claim registry are released; raw corpora are not redistributed.

1 Introduction

Language models read text as *tokens*, and a tokenizer's vocabulary determines how many tokens a given string becomes. Because sequence length drives context consumption and per-token accounting, a tokenizer that fragments a domain's recurring strings—resource identifiers, command flags, paths, configuration keys—into many pieces imposes a standing overhead on that domain. Subword tokenizers such as byte-pair encoding (BPE) [25], WordPiece [24], and Unigram/SentencePiece [12, 13] learn their inventories from a training corpus, so the corpus composition shapes which strings are cheap.

Technical documentation is a useful domain to study this effect: it mixes natural language with dense, repetitive technical syntax, and large, permissively licensed corpora exist. Azure documentation is a useful controlled case—sizable, public, and internally varied across services—so that we can hold out whole services for genuinely unseen evaluation. Crucially, in-domain gains must

*Author, affiliation, and any employer relationship are placeholders pending confirmation; see the repository's `publication/AUTHOR_INPUTS_REQUIRED.md`. This is an independent research artifact using only public documentation; it does not represent Microsoft's position and implies no endorsement.

be reported *with* out-of-domain regressions: a tokenizer specialized to one domain is expected to compress other text worse, and a practical deployment may prefer a mixed tokenizer that balances the two.

Existing production tokenizers (e.g., `tiktoken` [18]) differ from one another in training data, vocabulary size, normalization, and algorithm simultaneously, so comparing them cannot isolate the effect of the *training corpus* alone. We therefore run a controlled experiment: identical byte-level BPE tokenizers that differ only in their training corpus, trained on byte-matched budgets, with leakage-controlled held-out evaluation.

This paper measures *tokenization compression only*. It does not train or evaluate a language model, and makes no claim about model quality, accuracy, reasoning, inference latency, GPU throughput, or realized cost. Our contributions are:

1. A controlled corpus-composition study with matched byte budgets and fixed byte-level BPE configuration (General vs Azure vs Mixed, at 8K/16K/32K).
2. Cross-service and cross-repository held-out evaluation with explicit leakage controls (whole-service holdout, exact and near-duplicate removal, cross-corpus near-duplicate filtering).
3. An analysis of the specialization–generalization tradeoff via a 50/50 Mixed tokenizer.
4. Content-type, vocabulary-size, coverage, transfer, and CPU-runtime analyses, separating compression from runtime.
5. A reproducible artifact with source locks, raw per-document results, and a machine-verified claim registry in which every reported number is reconciled against its source file.

2 Related work

We organize related work by methodological relationship to this study rather than as a list of results.

Subword tokenization foundations. BPE for open-vocabulary neural machine translation [25], WordPiece [24], subword regularization and Unigram language-model tokenization [12], and SentencePiece [13] define the tokenizer families and the open-vocabulary motivation; byte-level BPE [20] and WordPiece-based encoders [8] are widely deployed. We do not introduce a new tokenizer algorithm; we vary the training corpus within a fixed byte-level BPE design.

Byte-level BPE and implementation baselines. Byte-level BPE [20] and its production implementations [11, 18] provide robust, full-coverage tokenization. We use one such implementation for the controlled tokenizers and treat production `tiktoken` encodings as observational baselines.

Domain-specific tokenizers and vocabularies. Prior work builds domain-specific vocabularies or tokenizers for scientific [4], legal [6], and legal/financial [5] corpora, and studies tokenizer choice for domains [7]; code and natural-language corpora similarly motivate technical text that mixes prose, identifiers, and snippets [9]. Notably, some domain-adapted models instead reuse a general-domain tokenizer [14], which is precisely the vocabulary-mismatch setting this line of work motivates. This study targets cloud/technical documentation and isolates *compression* rather than downstream model quality.

Vocabulary adaptation for existing models. A separate line adapts or augments an existing model's tokenizer—adaptive tokenization [23], adaptive BPE and vocabulary adaptation in fine-tuning [2, 10], general-to-specialized adaptation [16], and parameter-efficient vocabulary adaptation [3]. These methods modify an existing vocabulary and often evaluate downstream tasks; we instead train tokenizers from scratch on byte-matched corpora and report intrinsic compression only, so we make no model-substitution claim.

Pre-tokenization and scaling choices. Pre-tokenization [7, 13] and the amount of tokenizer training data [21] both affect the learned inventory. We fix byte-level BPE design choices and vary corpus composition and vocabulary size, observing diminishing sensitivity across 8K/16K/32K.

Fertility, compression, and fairness metrics. Token counts are corpus- and language-sensitive and can create inequities [19], motivating metrics beyond raw fertility [17] and studies of tokenizer choice [1, 22]. We adopt tokens per 1,000 UTF-8 bytes as the outcome and explicitly avoid downstream quality or fairness claims.

Leakage and evaluation controls. Deduplication improves the integrity of train/eval separation for language modeling [15]. We apply analogous exact and near-duplicate controls to tokenizer-compression evaluation.

Taken together, prior work establishes that tokenizer inventories are corpus-sensitive and can affect efficiency and downstream systems. Our contribution is a controlled, byte-matched, leakage-controlled case study measuring how that sensitivity manifests as *compression* on technical documentation, including the tradeoff between full specialization and mixed-domain training. This work sits primarily in the "trained from scratch on domain data" and "fertility/compression measurement" categories above.

3 Research questions and hypotheses

We pre-registered the hypotheses before inspecting the full results (repository file `research/EXPERIMENT_SPEC.md`).

- **RQ1.** Under matched conditions, does Azure-domain training reduce token count on unseen Azure services relative to general-English training?
- **RQ2.** What general-English regression accompanies full specialization?
- **RQ3.** Does a 50/50 Mixed corpus preserve most of the in-domain gain while limiting the regression?
- **RQ4.** Does the effect transfer to other technical-documentation domains?
- **RQ5.** Is the effect stable across 8K, 16K, and 32K vocabularies?
- **RQ6.** Does the effect persist on prose, rather than being explained only by code and markup?
- **RQ7.** How do compression and CPU tokenizer throughput relate?

Primary hypothesis. At a fixed byte-level BPE algorithm, vocabulary size, training-byte budget, normalization, sampling, and evaluation set, an Azure-trained tokenizer produces fewer tokens per 1,000 UTF-8 bytes on the Azure cross-service holdout than a byte-matched General tokenizer.

Table 1: Public sources, roles, retained Markdown file counts, and licenses. Sources are pinned by commit; raw text is not redistributed.

Source	Role	Files	License
MicrosoftDocs/azure-docs	train (anchor)	5,340	CC-BY-4.0
MicrosoftDocs/architecture-center	eval (transfer)	220	CC-BY-4.0
MicrosoftDocs/azure-devops-docs	eval (transfer)	220	CC-BY-4.0
MicrosoftDocs/sql-docs	eval (transfer)	220	CC-BY-4.0
MicrosoftDocs/fabric-docs	eval (transfer)	150	CC-BY-4.0
MicrosoftDocs/microsoft-365-docs	eval (transfer)	200	CC-BY-4.0
kubernetes/website	eval (external)	250	CC-BY-4.0
Salesforce/wikitext	train+eval	14,900	CC-BY-SA-3.0

Null hypothesis. After those controls, the Azure- and General-trained tokenizers do not differ in token count on the held-out set, i.e. the point estimate of the relative reduction is ≤ 0 or its 95% paired-bootstrap CI is compatible with no positive reduction.

4 Data

All sources are public and pinned by exact commit or dataset revision (repository file `data/manifests/sources.lock.` raw source text is not redistributed. Table 1 summarizes the sources; per-source licenses, attribution, and acquisition commands are in Appendix B.

Training corpora. Three byte-matched corpora of 33,554,432 UTF-8 bytes each: *General* (Wikipedia-derived English, WikiText-103-raw), *Azure* (20 service directories of `MicrosoftDocs/azure-docs`, 3,799 documents after deduplication and holdout), and *Mixed*, a deterministic document-by-document 50/50 interleave. Corpora are matched on cleaned bytes, not document count.

Evaluation corpora. Nine held-out sets: the Azure cross-service holdout (five whole services never seen in training, the *principal* set, 964 documents), an Azure same-service holdout, four Microsoft transfer sets (architecture, DevOps, data, broad-enterprise), one external technical control (Kubernetes website documentation), a held-out general-English set, and a diagnostic synthetic set.

Cleaning, classification, and leakage control. Documents are decoded as UTF-8, NFC-normalized (no lowercasing or accent stripping), stripped of front matter and build directives, segmented, and classified into content types by a rule-based classifier. Exact duplicates (SHA-256) and MinHash near-duplicates are removed from Azure before splitting; the General corpus is exact-deduplicated before the train/eval split (134 documents removed); and evaluation items that are near-duplicates of any training document are removed (9 items). Whole Azure services are held out so the principal evaluation is service-disjoint from training. We do not reproduce source-document text beyond minimal illustrative tokens.

5 Methodology

Controlled tokenizers (Track A). We train byte-level BPE tokenizers [20] with the Hugging Face `tokenizers` implementation [11], holding fixed: a byte-level pre-tokenizer and matching decoder, a full byte-level initial alphabet (so ordinary text never produces an unknown token), Unicode NFC

Table 2: Controlled experimental design. Only the training corpus is manipulated; all other factors are held fixed within a comparison.

Factor	Levels	Held constant
Training corpus (manipulated)	General, Azure, Mixed 50/50	—
Vocabulary size	8,192 / 16,384 / 32,768	within a comparison
Algorithm	byte-level BPE	implementation, pre-tokenizer, decoder
Training budget	33,554,432 UTF-8 bytes	byte-matched across corpora
Normalization	Unicode NFC	no lowercasing, no accent stripping
Seed	20260709	sampling, splitting, bootstrap

normalization without lowercasing or accent stripping, the special-token set, and a minimum merge frequency of two. The only variable is the training corpus (General, Azure, Mixed). Vocabulary sizes are 8,192, 16,384, and 32,768 (special tokens count toward the size), giving nine controlled tokenizers. Sampling, splitting, and the bootstrap are driven by a single seed (20,260,709). Corpora are sampled deterministically to exactly 33,554,432 bytes each; the Mixed corpus interleaves Azure and General documents 50/50 by bytes. Table 2 lists the design.

Observational baselines (Track B). We also encode the evaluation sets with production `tiktoken` encodings (`r50k_base`, `cl100k_base`, `o200k_base`) [18]. These differ from Track A and from one another in training data, vocabulary size, normalization, and intended model, so they are reported separately and are not a controlled comparison.

Primary metric. For a tokenizer T and evaluation corpus D,

$$\mathrm{tpk}(T, D) = 1000 \cdot \frac{\mathrm{tokens}(T, D)}{\mathrm{bytes}_{\text{UTF-8}}(D)}, \tag{1}$$

tokens per 1,000 UTF-8 bytes (lower is more compression), reported both *micro* (over all bytes) and *macro* (mean of per-document ratios). The primary effect size is the relative reduction of a specialized tokenizer A versus the matched General tokenizer G,

$$\Delta(A \text{ vs } G, D) = 100 \cdot \frac{\mathrm{tpk}(G, D) - \mathrm{tpk}(A, D)}{\mathrm{tpk}(G, D)}, \tag{2}$$

where *positive* Δ means the specialized tokenizer uses *fewer* tokens. Secondary metrics include characters/bytes per token, tokens per whitespace word, identifier fragmentation, vocabulary utilization, exact round-trip and unknown-token rates, and CPU throughput.

Statistics. Token counts are deterministic; uncertainty comes from the sampled documents. We use a paired bootstrap over documents (10,000 resamples; the same documents resampled for both tokenizers), reporting 95% confidence intervals for Δ, a document-level win rate, and stratification by evaluation set, content type, and repository. The primary hypothesis (Section 3) is the single confirmatory comparison; the transfer, vocabulary-size, content-type, and per-repository breakdowns are secondary and exploratory, and their intervals are not corrected for multiple comparisons, so we read them as descriptive rather than confirmatory.

Table 3: Main results at 16K: token reduction (%) relative to the byte-matched General tokenizer on the Azure cross-service holdout and on held-out general English. Positive = fewer tokens.

Tokenizer (16K)	Azure cross-service reduction (%)	General English reduction (%)
Azure	+36.68	−34.75
Mixed (50/50)	+34.21	−4.39
General	0 (ref.)	0 (ref.)

Azure primary 95% CI [35.77, 37.62], win rate 1.0, $n = 964$.
Positive = specialized tokenizer uses fewer tokens per 1,000 UTF-8 bytes.

Runtime. CPU tokenizer throughput is measured separately from compression, with a fixed thread count, warm-up, five measured rounds, randomized tokenizer order per round, identical documents and batches, and serialized execution, on an Apple M3 Pro (12 cores, macOS). This is CPU tokenizer throughput, not model-serving throughput.

Verification. An automated gate re-checks byte-matching, leakage controls, exact round-trip and unknown-token counts, totals reconciliation, provenance completeness, artifact-distribution policy, and source pinning without retraining. Every number in this paper is additionally reconciled against its source file by a claim registry (Appendix E).

6 Results

We report the 16K slice unless stated. Positive relative reduction means the specialized tokenizer uses fewer tokens. All values are reconciled against source files by the claim registry (Appendix E).

Primary result. On the Azure cross-service holdout (five whole held-out services, $n = 964$ documents), the Azure tokenizer uses 240.18 tokens per 1,000 UTF-8 bytes versus 379.30 for the byte-matched General tokenizer, a relative reduction of 36.68% (95% paired-bootstrap CI [35.77%, 37.62%]; document win rate 1.000). Table 3 and Figure 1 summarize the result and its stability across vocabulary sizes.

General-English regression (RQ2). Full specialization has a cost: on held-out general English the Azure tokenizer is −34.75% relative to General (95% CI [−35.10%, −34.40%])—it uses more tokens on out-of-domain text.

Mixed-tokenizer tradeoff (RQ3). The 50/50 Mixed tokenizer keeps most of the in-domain gain (+34.21% on the Azure holdout) while nearly eliminating the out-of-domain penalty (−4.39% on general English) (Figure 2).

Transfer (RQ4). The advantage decays with domain distance and persists on non-Microsoft documentation: +35.03% (Azure same-service), +32.94% (Microsoft data), +31.76% (Microsoft DevOps), +31.18% (Kubernetes, external), +26.03% (Microsoft architecture), and +25.89% (Microsoft broad-enterprise); all confidence intervals exclude zero. Table 4 and Figure 3 report transfer with the general-English regression.

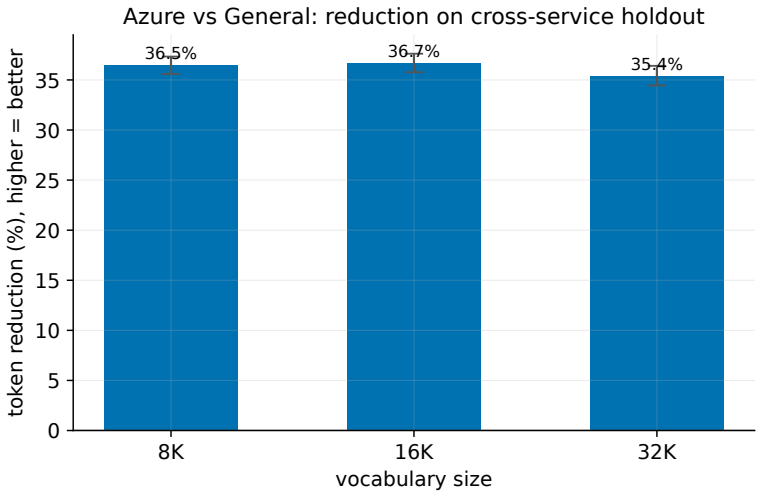


Figure 1: Relative token reduction (%) of the Azure vs byte-matched General tokenizer on the Azure cross-service holdout, by vocabulary size, with 95% paired-bootstrap confidence intervals. Higher is better.

Table 4: Transfer of the Azure vs General reduction (%, 16K) across evaluation sets, plus vocabulary-size sensitivity on the principal holdout. Positive = fewer tokens.

Evaluation set	**Azure vs General reduction (%)**
Azure cross-service holdout (principal)	+36.68
Azure same-service holdout	+35.03
Microsoft data docs	+32.94
Microsoft DevOps docs	+31.76
Kubernetes (external technical)	+31.18
Microsoft architecture docs	+26.03
Microsoft broad-enterprise docs	+25.89
General English	−34.75
Vocabulary sensitivity (principal): 8K +36.46, 16K +36.68, 32K +35.41.	

Vocabulary-size sensitivity (RQ5). The reduction on the principal holdout is stable across the tested range: +36.46% (8K), +36.68% (16K), +35.41% (32K) (Figure 4), suggesting corpus composition mattered more than the tested vocabulary range.

Content type (RQ6). By dominant content type on the principal holdout: prose-dominant documents ($n = 900$) show +36.75% (95% CI [35.87%, 37.63%]), headings +37.49%, tables +44.30%, and inline code +33.28% (Figure 5). Because the largest stratum (prose) still shows a large reduction, the effect is not explained only by code fences or configuration syntax.

Coverage and runtime (RQ7). Across 12 encoders × 9 evaluation sets (59,160 rows), reconstruction is exact (100% round-trip) with zero unknown tokens. Median CPU tokenizer throughput is 5.1–5.5 MB/s for the controlled tokenizers and 11.6–19.8 MB/s for the `tiktoken` baselines (Figure 6); the controlled tokenizers compress more but run slower on CPU than the production-optimized baselines. Compression and CPU throughput are distinct dimensions, and CPU throughput is not

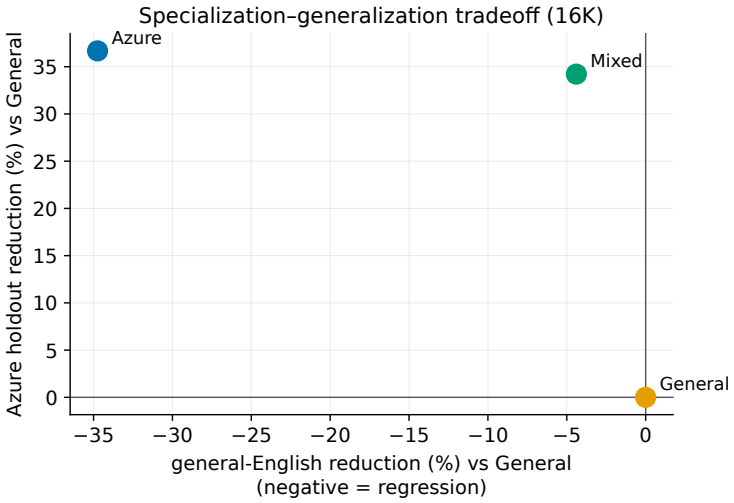


Figure 2: Specialization–generalization tradeoff at 16K: token reduction (%) relative to General on the Azure holdout (y) versus held-out general English (x). Up and to the right is better on both axes.

Table 5: Coverage and CPU-runtime summary. CPU tokenizer throughput is not model-serving throughput.

Measure	**Value**
Evaluation rows	59,160 (12 encoders × 9 sets)
Exact round-trip	100%
Unknown tokens (`[UNK]`)	0
CPU throughput, Track A (MB/s)	5.13–5.47
CPU throughput, Track B `tiktoken` (MB/s)	11.61–19.83
Automated verification gate	8/8

model-serving throughput. Table 5 summarizes coverage and runtime.

7 Analysis and discussion

The results are consistent with the Azure tokenizer allocating vocabulary to domain-recurring strings and documentation patterns, so that resource-provider names, command flags, paths, and configuration keys are encoded in fewer pieces. The mechanism appears in fragmentation and vocabulary-utilization measurements in the repository: on the principal holdout the Azure tokenizer encodes resource-provider identifiers, URLs, command flags, and camel/snake identifiers in fewer tokens than the General tokenizer (for example, roughly half as many tokens per resource-provider identifier in the repository's fragmentation table), and it uses a larger fraction of its vocabulary on Azure text.

Several observations bound the interpretation. First, the effect *persists on prose* (the largest content stratum, $n = 900$, still shows +36.75%), which weakens the alternative that gains come only from code fences or configuration syntax. Second, the advantage *transfers* to non-Microsoft Kubernetes documentation (+31.18%), suggesting overlap in general technical-documentation style rather than only Azure-specific terminology. Third, the reduction is *stable across* 8K/16K/32K,

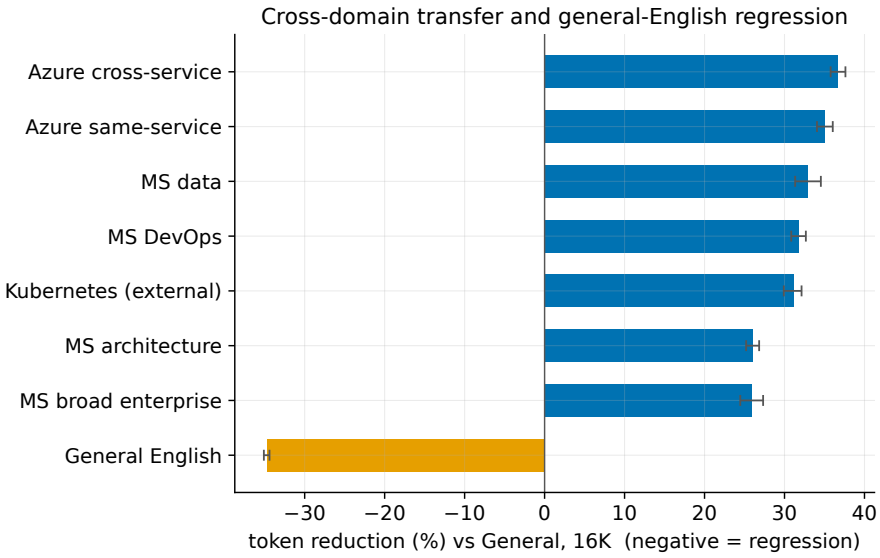


Figure 3: Azure vs General reduction (%, 16K) across evaluation sets with 95% paired-bootstrap CIs. The general-English bar is negative (a regression).

suggesting corpus composition mattered more than the tested vocabulary range. Fourth, the Mixed tokenizer is the central practical result: it recovers most of the in-domain gain with a small out-of-domain penalty, which matters when a single tokenizer must serve mixed traffic.

We are explicit about a design limit: comparing Azure documentation with general English changes *more than subject domain*. It also changes technical writing style, markup distribution, identifier density, and document structure. The transfer and prose-only analyses narrow, but do not fully isolate, Azure-specific vocabulary from general technical-documentation specialization (Section 8). Finally, faster compression does not imply faster CPU tokenization—the controlled tokenizers compress more yet run slower on CPU than the production baselines—and none of these measurements support any downstream model claim.

What this study is and is not

This study is: a tokenizer-training and *compression* benchmark; a controlled comparison of corpus composition; a held-out technical-documentation evaluation; a study of domain/general tradeoffs; a CPU-only reproducible artifact.

This study is not: an LLM quality evaluation; a model-adaptation method; a serving or GPU benchmark; a billing study; an endorsement of replacing production tokenizers; or a claim that fewer tokens always means a better system.

8 Limitations

We state limitations plainly and do not minimize them.

Training budget and design. Corpora are byte-matched at a small 32 MiB budget, bounded by the available cleaned Azure training text; absolute token counts are budget-dependent. The controlled

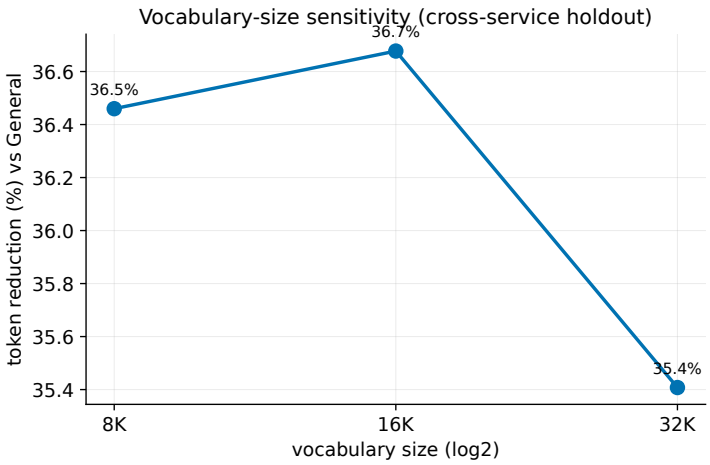


Figure 4: Vocabulary-size sensitivity: Azure vs General reduction (%) on the cross-service holdout at 8K/16K/32K.

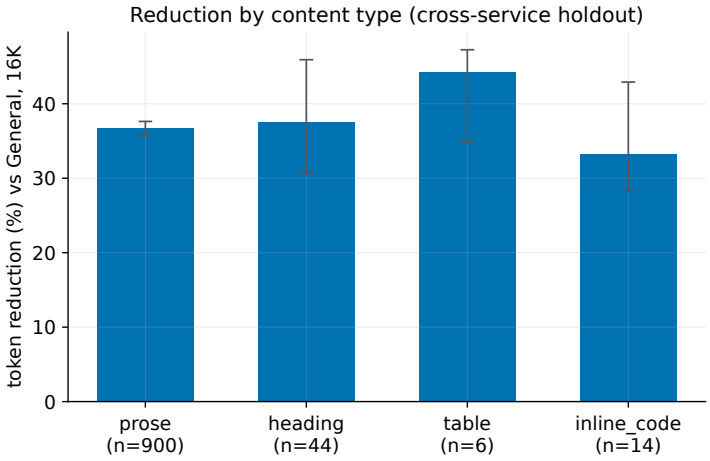


Figure 5: Reduction by dominant content type on the cross-service holdout (16K), with 95% CIs and per-stratum document counts.

matrix uses a single tokenizer algorithm (byte-level BPE). Sources are technical documentation, so results may not extend to chat, source-code repositories, prose, or all enterprise text.

The Azure-versus-general-English confound (most important). Comparing Azure documentation with general English changes more than subject domain; it also changes technical writing style, markup distribution, identifier density, and document structure. The transfer and prose-only analyses narrow this concern but do *not* fully isolate Azure-specific vocabulary from general technical-documentation specialization. A matched general-technical documentation training control is the priority follow-up (Section 10).

Scope of measurement. We train no language model and evaluate no downstream task; there is no GPU or serving benchmark. CPU runtime was measured on a thermally variable laptop, so small runtime differences are not overinterpreted, and CPU tokenizer throughput is not model-serving

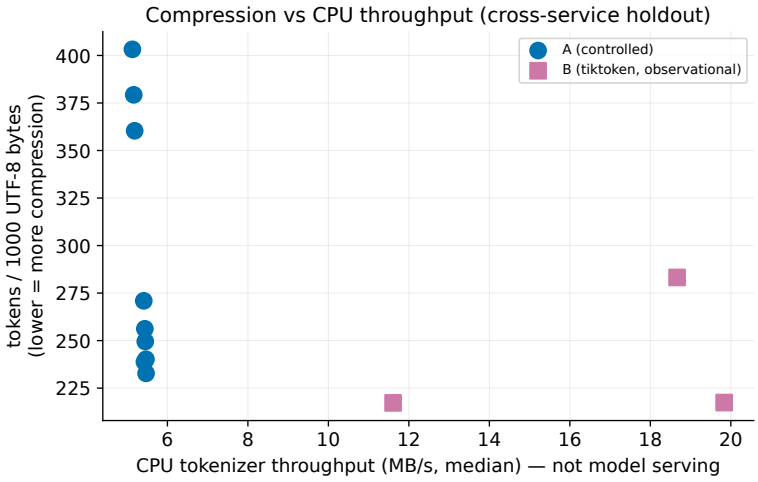


Figure 6: Compression (tokens per 1,000 UTF-8 bytes, lower is better) versus measured CPU tokenizer throughput on the cross-service holdout. Track A are the controlled tokenizers; Track B (`tiktoken`) are observational baselines.

throughput.

Data and comparison caveats. Results are tied to pinned public-source snapshots that change over time; raw corpora are not redistributed, and reproduction downloads the pinned sources. Existing production tokenizers (Track B) are observational, not controlled. Any context-window or token-billing figure elsewhere is a derived estimate, not a measured GPU or billing outcome. We claim no universal domain-tokenizer superiority.

Affiliation. If the author is affiliated with the vendor of the studied documentation, that relationship is disclosed where applicable; this study uses only public documentation, is not an official result of that vendor, and implies no endorsement.

9 Ethics, licensing, and broader impact

Sources and licensing. All data are public documentation. Microsoft documentation sources are used under CC-BY-4.0 (documentation) with MIT code samples; the Kubernetes website content is CC-BY-4.0; the general-English corpus (WikiText-103) is Wikipedia-derived under CC-BY-SA. Sources are pinned by commit and attributed in the repository (`licenses/`). No user data and no private content are used.

Redistribution boundaries. We do not republish raw source corpora; reproduction downloads the pinned public sources. We release scripts, source manifests, content hashes, per-document result rows, processed tables, and figures. The Azure-only tokenizer artifact is redistributable with attribution; the General and Mixed tokenizers are derived in part from WikiText and are kept local-use-only pending clarification of the applicable share-alike terms.

Broader impact. Token-count differences across domains can translate into uneven context capacity and per-token accounting across users and domains; studying compression makes such

disparities measurable. We caution against over-reading compression as model quality: fewer tokens is an efficiency property, not evidence of better understanding, accuracy, or safety.

Environmental scope. The experiment is CPU-only and trains no language model; its compute footprint is small (tokenizer training and evaluation complete in minutes on a laptop).

Non-endorsement and review. This is independent research using public documentation and does not represent the position of any vendor. We make no IRB claim (no human-subjects data). Author affiliation and any required employer publication review are recorded as open items in the repository.

10 Future work

The priority follow-up is a *matched general-technical documentation* training control that shares technical style, markup, and identifier density with the Azure corpus but not its Azure-specific vocabulary, to separate general technical-documentation specialization from Azure specialization. Beyond that: (i) more corpus mixture ratios between the General and Azure extremes; (ii) larger training-data budgets to test whether the reduction changes with scale; (iii) other tokenizer algorithms (e.g., Unigram, WordPiece) under matched controls; (iv) deeper vocabulary-utilization and token-category allocation analysis; and (v) temporal holdouts across documentation snapshots.

Separately, and clearly labeled as future GPU/model work not attempted here, one could integrate an adapted tokenizer into a small model trained from scratch to measure any downstream effect, and measure real prefill, KV-cache, latency, and throughput with a compatible model. Evaluating code-heavy and multilingual technical documentation, and vocabulary-replacement or adaptation methods rather than full retraining, are additional directions. None of these downstream or serving results are claimed by the present study.

11 Conclusion

Under a controlled, byte-matched comparison at a matched 16K vocabulary, training a byte-level BPE tokenizer on Azure documentation reduced token count on an unseen cross-service Azure holdout by 36.68% relative to a general-English tokenizer (95% CI $[35.77\%, 37.62\%]$), with a general-English regression of -34.75%. A 50/50 Mixed tokenizer retained most of the in-domain gain ($+34.21\%$) with a much smaller general-English penalty (-4.39%), and the advantage transferred partially to other Microsoft and external technical-documentation corpora. These findings concern *tokenization compression only*, on the evaluated corpora and under the stated controls; they imply nothing about model output quality, inference latency, GPU throughput, or cost. The controlled comparison isolates the training corpus, whereas existing production tokenizers differ along multiple axes and serve only as observational context. All numbers are reproducible from a released, machine-verified artifact.

Reproducibility statement

All results regenerate from a public, version-controlled repository. Every empirical result value reported in this paper (effect sizes, confidence intervals, coverage counts, and runtime summaries) is registered in a machine-readable claim registry (`publication/claims.yaml`) and re-checked against its authoritative source file by `make verify-claims`; experimental constants (seed, budgets,

batch sizes, vocabulary sizes) are recorded in the repository configs. The analysis-only path (`make verify`, `make verify-claims`, `make figures`, `make report`) runs from committed result files without downloading any corpus; full reproduction downloads the pinned public sources. See Appendix F.

Data and code availability

Code, source manifests with pinned commits, content hashes, per-document result rows, processed tables, and figures are released under the repository's terms. Raw source corpora are not redistributed; reproduction downloads the pinned public sources. The Azure-only tokenizer artifact is distributable with attribution; the General and Mixed (WikiText-derived) tokenizers are local-use-only pending license review.

References

[1] Mehdi Ali, Michael Fromm, Klaudia Thellmann, Richard Rutmann, Max Lübbering, Johannes Leveling, Katrin Klug, Jan Ebert, Niclas Doll, Jasper Schulze Buschhoff, Charvi Jain, Alexander Arno Weber, Lena Jurkschat, Hammam Abdelwahab, Chelsea John, Pedro Ortiz Suárez, Malte Ostendorff, Samuel Weinbach, Rafet Sifa, Stefan Kesselheim, and Nicolas Flores-Herr. Tokenizer choice for LLM training: Negligible or crucial? In *Findings of the Association for Computational Linguistics: NAACL 2024*, 2024. URL https://arxiv.org/abs/2310.08754.

[2] Gunjan Balde, Soumyadeep Roy, Mainack Mondal, and Niloy Ganguly. Adaptive BPE tokenization for enhanced vocabulary adaptation in finetuning pretrained language models. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 14724–14733, Miami, Florida, USA, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-emnlp.863. URL https://aclanthology.org/2024.findings-emnlp.863/.

[3] Gunjan Balde, Soumyadeep Roy, Mainack Mondal, and Niloy Ganguly. Learning faster with better tokens: Parameter-efficient vocabulary adaptation for specialized text summarization, 2026. URL https://arxiv.org/abs/2605.17379. Accepted to ACL 2026.

[4] Iz Beltagy, Kyle Lo, and Arman Cohan. SciBERT: A pretrained language model for scientific text. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, 2019. URL https://arxiv.org/abs/1903.10676.

[5] Michael J. Bommarito, Daniel Martin Katz, and Jillian Bommarito. KL3M tokenizers: A family of domain-specific and character-level tokenizers for legal, financial, and preprocessing applications, 2025. URL https://arxiv.org/abs/2503.17247.

[6] Ilias Chalkidis, Manos Fergadiotis, Prodromos Malakasiotis, Nikolaos Aletras, and Ion Androutsopoulos. LEGAL-BERT: The muppets straight out of law school. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, 2020. URL https://arxiv.org/abs/2010.02559.

[7] Gautier Dagan, Gabriel Synnaeve, and Baptiste Rozière. Getting the most out of your tokenizer for pre-training and domain adaptation. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 9784–9805. PMLR, 2024. URL https://proceedings.mlr.press/v235/dagan24a.html.

[8] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 2019. URL https://arxiv.org/abs/1810.04805.

[9] Zhangyin Feng, Daya Guo, Duyu Tang, Nan Duan, Xiaocheng Feng, Ming Gong, Linjun Shou, Bing Qin, Ting Liu, Daxin Jiang, and Ming Zhou. CodeBERT: A pre-trained model for programming and natural languages. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, 2020. URL https://arxiv.org/abs/2002.08155.

[10] Christian Herold, Michael Kozielski, Nicholas Santavas, Yannick Versley, and Shahram Khadivi. Vocabulary customization for efficient domain-specific LLM deployment, 2025. URL https://arxiv.org/abs/2509.26124. Accepted at NEURIPS 2025 CCFM Workshop.

[11] Hugging Face. tokenizers. GitHub repository, 2026. URL https://github.com/huggingface/tokenizers. Accessed 2026-07-10.

[12] Taku Kudo. Subword regularization: Improving neural network translation models with multiple subword candidates. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 66–75, Melbourne, Australia, 2018. Association for Computational Linguistics. doi: 10.18653/v1/P18-1007. URL https://aclanthology.org/P18-1007/.

[13] Taku Kudo and John Richardson. Sentencepiece: A simple and language independent subword tokenizer and detokenizer for neural text processing. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 66–71, Brussels, Belgium, 2018. Association for Computational Linguistics. doi: 10.18653/v1/D18-2012. URL https://aclanthology.org/D18-2012/.

[14] Jinhyuk Lee, Wonjin Yoon, Sungdong Kim, Donghyeon Kim, Sunkyu Kim, Chan Ho So, and Jaewoo Kang. BioBERT: A pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 2019. doi: 10.1093/bioinformatics/btz682. URL https://arxiv.org/abs/1901.08746.

[15] Katherine Lee, Daphne Ippolito, Andrew Nystrom, Chiyuan Zhang, Douglas Eck, Chris Callison-Burch, and Nicholas Carlini. Deduplicating training data makes language models better. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8424–8445. Association for Computational Linguistics, 2022. doi: 10.18653/v1/2022.acl-long.577. URL https://aclanthology.org/2022.acl-long.577/.

[16] Vijini Liyanage and François Yvon. AdaptBPE: From general purpose to specialized tokenizers, 2026. URL https://arxiv.org/abs/2601.21665. EACL 2026.

[17] Mir Tafseer Nayeem, Sawsan Alqahtani, Md Tahmid Rahman Laskar, Tasnim Mohiuddin, and M Saiful Bari. Beyond fertility: Analyzing STRR as a metric for multilingual tokenization evaluation, 2025. URL https://arxiv.org/abs/2510.09947. NeurIPS 2025 Workshop.

[18] OpenAI. tiktoken. GitHub repository, 2026. URL https://github.com/openai/tiktoken. Accessed 2026-07-10.

[19] Aleksandar Petrov, Emanuele La Malfa, Philip H. S. Torr, and Adel Bibi. Language model tokenizers introduce unfairness between languages. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL https://arxiv.org/abs/2305.15425.

[20] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. Technical report, OpenAI, 2019. URL https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf.

[21] Varshini Reddy, Craig W. Schmidt, Yuval Pinter, and Chris Tanner. How much is enough? the diminishing returns of tokenization training data, 2025. URL https://arxiv.org/abs/2502.20273.

[22] Phillip Rust, Jonas Pfeiffer, Ivan Vulić, Sebastian Ruder, and Iryna Gurevych. How good is your tokenizer? on the monolingual performance of multilingual language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 3118–3135. Association for Computational Linguistics, 2021. doi: 10.18653/v1/2021.acl-long.243. URL https://aclanthology.org/2021.acl-long.243/.

[23] Vin Sachidananda, Jason Kessler, and Yi-An Lai. Efficient domain adaptation of language models via adaptive tokenization. In *Proceedings of the Second Workshop on Simple and Efficient Natural Language Processing*, pages 155–165, Virtual, 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.sustainlp-1.16. URL https://aclanthology.org/2021.sustainlp-1.16/.

[24] Mike Schuster and Kaisuke Nakajima. Japanese and korean voice search. In *2012 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 5149–5152. IEEE, 2012. doi: 10.1109/ICASSP.2012.6289079. URL https://doi.org/10.1109/ICASSP.2012.6289079.

[25] Rico Sennrich, Barry Haddow, and Alexandra Birch. Neural machine translation of rare words with subword units. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1715–1725, Berlin, Germany, 2016. Association for Computational Linguistics. doi: 10.18653/v1/P16-1162. URL https://aclanthology.org/P16-1162/.

A Experimental matrix

The controlled matrix comprises nine tokenizers: General, Azure, and Mixed (50/50) training corpora crossed with vocabulary sizes 8,192, 16,384, and 32,768. Within each vocabulary size, the three tokenizers are identical except for the training corpus; each is trained on exactly 33,554,432 UTF-8 bytes. Track B adds three observational `tiktoken` encodings, giving 12 encoders evaluated on 9 held-out sets. Per-tokenizer provenance (configuration hash, artifact hash, byte counts, merge count, package versions, machine metadata, and redistribution flag) is recorded in `artifacts/tokenizers/*.json`.

B Dataset details

Table 1 lists sources, roles, retained Markdown file counts, and licenses. Each source is pinned by exact commit or dataset revision in `data/manifests/sources.lock.json`; per-source acquisition commands, selection method (git sparse checkout, GitHub API blob download, or Hugging Face parquet), attribution, and redistribution status are in `docs/datasets-and-licenses.md` and

`licenses/`. Raw source text is not redistributed. Microsoft documentation sources are CC-BY-4.0; the Kubernetes website is CC-BY-4.0; WikiText-103 is Wikipedia-derived under CC-BY-SA and is used local-use-only.

C Additional results

The repository reports secondary analyses that support the main text: identifier-fragmentation distributions (mean tokens per resource-provider identifier, URL/path, command flag, and camel/snake identifier), vocabulary utilization (fraction of vocabulary used on the principal holdout and occurrence concentration), and both micro and macro aggregations of the primary metric. These tables are generated from the committed per-document rows (`results/raw/eval_rows.parquet`) and are listed in `results/DATA_DICTIONARY.md`. All confidence intervals are 95% paired-bootstrap intervals over documents (10,000 resamples), paired at the document level and stratified by evaluation set, content type, and repository.

D Runtime details

CPU tokenizer throughput was measured with a fixed thread count, one warm-up round, five measured rounds, tokenizer order randomized per round, identical documents and batches (batch sizes $\{1, 8, 32, 128\}$; short/medium/long buckets), and serialized execution on an Apple M3 Pro (12 cores, macOS). Median throughput was 5.1–5.5 MB/s for the controlled tokenizers and 11.6–19.8 MB/s for the `tiktoken` baselines. Runtime was measured on a thermally variable laptop; small differences are not overinterpreted. This is CPU tokenizer throughput, not model-serving throughput, and it is reported separately from compression.

E Verification and claim registry

Every empirical result value in this paper (effect sizes, confidence intervals, coverage counts, and runtime summaries) is registered in a machine-readable claim registry (`publication/claims.yaml`); experimental constants are recorded in the repository configs. For each registered claim the registry records an extraction specification (source file and column/key), a frozen display value, and a tolerance; `make verify-claims` re-extracts each value from its authoritative result file and fails on any mismatch. An automated experiment gate (`make verify`) independently re-checks byte-matching of corpora, leakage controls, exact round-trip and unknown-token counts (100% round-trip, zero unknown tokens across 59,160 rows), totals reconciliation, provenance completeness, artifact-distribution policy, and source pinning, without retraining. The independently recomputed primary micro reduction from the per-document rows is 36.677%, matching the reported 36.68%.

F Reproducibility

The environment is pinned with `uv` (Python 3.12; `uv.lock`). Two paths are supported. *Analysis-only* reproduction runs from committed result files without downloading any corpus:

```
uv sync; make verify; make verify-claims; make figures; make report
```

Full reproduction downloads the pinned public sources and recomputes:

```
make acquire-full audit-full train-all evaluate-all;
PROFILE=full make analyze report figures verify
```

Full reproduction takes roughly 20–25 minutes on an Apple M3 Pro and a transient ∼2.7 GB of disk (a git-ignored source clone plus working data); the analysis-only path runs in minutes with no downloads. All stochastic steps are driven by a single seed (20,260,709); the main known nondeterminism is CPU runtime timing on a thermally variable laptop.