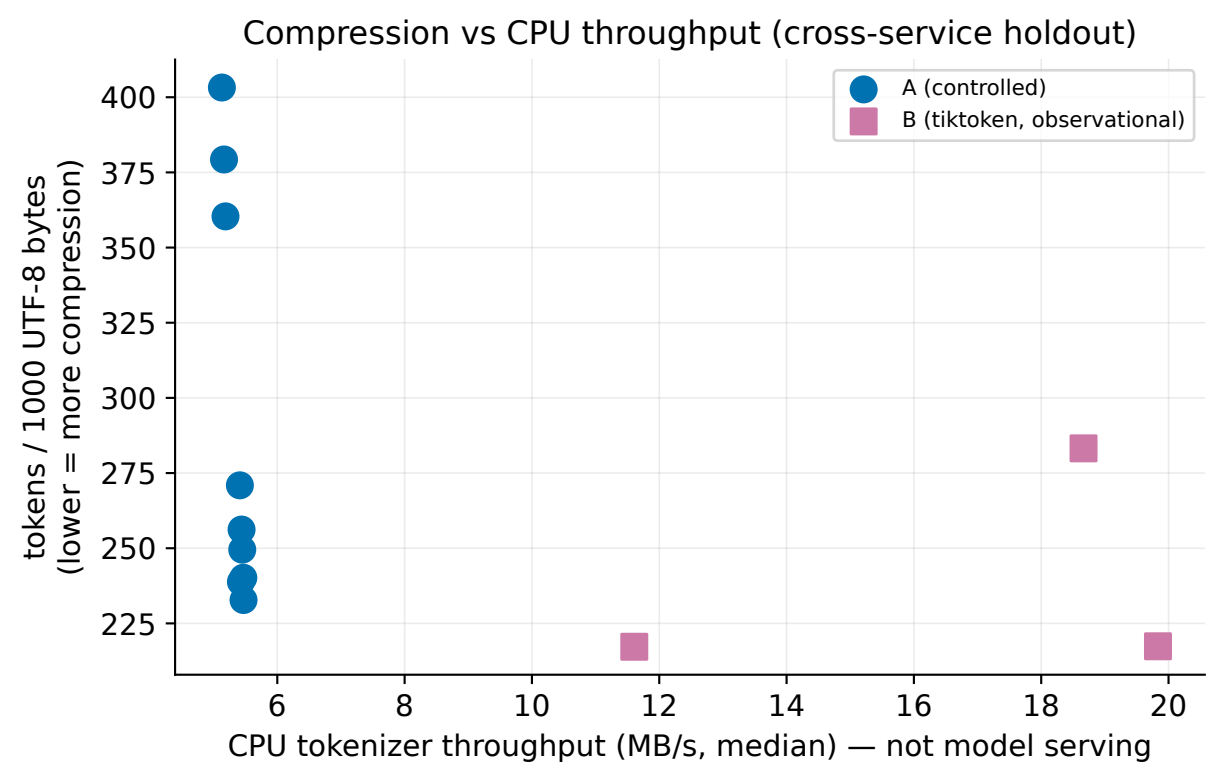
Compression vs CPU throughput (cross-service holdout)
A (controlled)
B (tiktoken, observational)
tokens / 1000 UTF-8 bytes
(lower = more compression)
400
375
350
325
300
275
250
225
6
8
10
12
14
16
18
20
CPU tokenizer throughput (MB/s, median) — not model serving